

IMPACT STATEMENT: INTELLIGENT CONSUMER COMPLAINT CLASSIFICATION

BY TEAM UPEMBAPARK

Our project explores the potential societal impact of using a fine-tuned transformer model to classify consumer complaints. The intended benefit is to support organisations in processing large volumes of complaints more efficiently and consistently, potentially enabling recurring consumer issues to be identified more quickly and helping complaints reach the appropriate area for further review.

However, automated complaint classification also presents ethical and social risks. Consumer complaints can contain sensitive personal information, and inappropriate handling of this data could create privacy risks. The model may also perform differently across complaints that use different language, levels of detail, or ways of describing an experience. Incorrect or biased classifications could result in some complaints being deprioritised or directed to the wrong category, potentially affecting consumers who are already vulnerable or less able to navigate formal complaint processes.

For this reason, UpembaPark's goal considers fairness, dignity, autonomy, and community well-being alongside model performance. The system is intended to assist rather than replace human decision-making. Human oversight would remain important when classifications could affect how a complaint is handled, particularly where the model is uncertain or where complaints involve sensitive circumstances.

Over time, the technology could contribute to more efficient and data-informed complaint management, but there is also a risk of organisations becoming overly dependent on automated classifications. Our assessment therefore recognises that the project's impact cannot be determined by accuracy alone. Responsible deployment would require continued evaluation for bias, protection of consumer information, monitoring of errors, and mechanisms for human review and correction. Because our model is being developed and evaluated in a project setting, its real-world effects remain uncertain and should not be assumed without further testing in representative deployment environments.